Lesson Plan

Using Affect Recognition to Establish Good Research Habits

Course: AI Literacy — College of Arts and Sciences Core Curriculum
Level: Sophomore
Session Length: Two class sessions × 80 minutes = approximately 150 minutes of instructional time
Lesson Title: Using Affect Recognition to Establish Good Research Habits

---

Overview

This lesson uses affect recognition technology as a concrete, socially relevant case study through which students develop transferable research skills. Rather than treating AI literacy as abstract digital citizenship, the lesson grounds it in real companies, real sources of varying types, and a real ethical debate — asking students to practice the habits of mind that distinguish careful inquiry from passive consumption of information. The primary text is an AI-generated research document on affect recognition, which introduces an additional layer of critical examination: students are evaluating not only the sources cited, but the nature of the document itself.

The lesson runs across two class sessions. Session One covers Task 0 as an out-of-class warm-up, then moves through Tasks 1, 2, and 3 in class, closing with the four structured reflection questions. Session Two opens with a 20-minute group discussion of what Session One revealed, then moves into Tasks 4 and 5.

Learning Outcomes

By the end of this lesson, students will be able to:

1. Apply source literacy skills to evaluate the credibility, purpose, and potential bias of both academic and industry-produced texts.

Students will practice tracing the origins of claims — identifying who wrote a source, what data they used, where that data came from, and what institutional or financial interests may have shaped their conclusions.

2. Demonstrate industry literacy by analyzing how technology companies communicate — and strategically omit — information about the products they build and deploy.

Students will read corporate About pages as rhetorical documents, examining what companies claim about privacy, consent, and public benefit alongside what they do not say.

3. Evaluate the relationship between scientific knowledge, commercial application, and social consequence in the context of AI systems.

Students will connect the contested empirical foundations of affect recognition to its real-world deployment, recognizing how technologies built on disputed science can cause material harm — particularly to populations whose expression, culture, or neurology falls outside the norms encoded in training data.

Materials

- Claude-authored research document: *Focus on Affect Recognition Technology* — available on Brightspace for reference, but students are **not** expected or asked to read it before class. It will be distributed at the start of Task 1 in Session One. This is intentional: students should arrive with Crawford's arguments fresh in mind and their own assumptions intact, without prior exposure to an AI-generated synthesis that could shape their reading before the lesson begins.

- Kate Crawford, *Atlas of AI*, Chapter 5: "Affect" (assigned reading, completed before the pre-class quiz)
- Access to the internet (for Tasks 2 and 3)
- In-class assignment sheet (Tasks 0–5, distributed at the start of Session One)
- Brightspace submission portal for written responses

Before Session One: Pre-Class Quiz and Task 0 (Brightspace)

Pre-Class Quiz

Students complete the five-question pre-class quiz before arriving to Session One. This confirms engagement with Crawford's Chapter 5, surfaces prior assumptions about emotional expression, and primes students for the meta-quality of working with an AI-authored document. Instructor reviews aggregate Likert responses (Question 2) before Session One to calibrate the opening discussion.

Task 0 (Out of Class — not counted in session time)

Students complete Task 0 on Brightspace before arriving to Session One. The prompt is:

> *"Before you engage with today's document, write two sentences about what you expect from a piece of research written by an AI. What assumptions are you bringing to it?"*

Students should save their response. They will return to it at the end of Session Two. Estimated time: 5 minutes.

Session One (80 minutes)

Opening Discussion — 5 minutes

Open by displaying the aggregate results from Question 2 of the pre-quiz without individual attribution. Ask the room: *"Some of you came in with positive or neutral feelings about affect recognition. Some came in skeptical. What do you think accounts for that range?"*

Keep this brief. The goal is simply to activate prior knowledge and orient the class — not to resolve anything. The deeper conversation happens in Session Two once students have done the investigative work.

Close the opening by returning to Question 4 of the pre-quiz — the True/False item about emotional expression and genetics. Ask for a quick show of hands on how students answered. Do not reveal the "correct" framing. Tell students they will return to this question at the end of Session Two.

Task 1: Read the Document — 20 minutes

The Claude-authored document *Focus on Affect Recognition Technology* is distributed now, at the start of Task 1. This is the first time students encounter it.

The instructor will call on eight students to read sections of the document aloud to the group, approximately two minutes of reading each, for a total of roughly 16 minutes of read-aloud time. This pacing is intentional — the document covers dense material across several domains including neurodevelopmental psychology, EU regulatory law, insurance underwriting, and AI ethics, and reading aloud together ensures the whole class moves through it at a pace where the content can actually land. It also establishes a shared reference point for the investigative work that follows.

The remaining four minutes allow the instructor to ask one or two brief orienting questions before students move into Tasks 2 and 3 — for example: *"What stood out to you? What do you want to look at more closely?"*

Task 2: Company Research — 15 minutes

Students select two companies or research architectures mentioned in the document, locate their websites, find the About section, and answer the following four questions for each company:

- **Q1:** Is the company for-profit?
- **Q2:** How is the company profiting from the technology?
- **Q3:** Does the company make any statements about protecting public privacy?
- **Q4:** Is the technology being used without informed consent? Note: fine print buried in lengthy consent disclosures does not constitute informed consent.

Task 3: Source Analysis — 25 minutes

Students select **two** sources cited in the document — these do not need to be academic sources; any cited source may be chosen — follow the provided links, and conduct a structured critical analysis of each. For each source, students identify and respond to the following:

- Who wrote the article or research?
- What was the author's stated purpose? If no purpose was stated, was one implied? If so, what was it?
- Did you fact-check any of the claims made? Did the authors use their own data, or did they cite data from another source? From where did that data come, and how was it collected?
- What were the main points of the article or source?
- Did you uncover any potential bias? For example, was the author also part of a company or organization that benefits from their conclusions?
- What assumptions, if any, did the author make to support their claims? Is there reason to question any of those assumptions?
- Rate the source on a scale of 1 to 5, where 1 = highly biased or reliant on unsupported assumptions, and 5 = low bias with claims backed by evidence and fact.

Instructor circulates throughout Tasks 2 and 3. Students who complete Task 2 quickly should invest the additional time in Task 3, which is the more analytically demanding of the two.

Reflection on Tasks 2 and 3 — 10 minutes

Students complete the following four written reflection prompts individually before Session One ends. These are submitted as part of the assignment.

R1. Looking at the two company About pages you examined in Task 2: What did the companies say about privacy and public benefit? What did they *not* say? In your view, does the absence of certain information tell you something? Explain your thinking.

R2. In Task 3, you evaluated two sources for potential bias and unsupported assumptions. Did either source change your expectations — either because it was more rigorous than you anticipated, or less? What specific feature of the source shaped that judgment?

R3. After doing Tasks 2 and 3, do you read the claims in the AI-generated document differently than you did after Task 1? Is there anything in the document you now trust more, trust less, or want to investigate further? Be specific.

R4. Tasks 2 and 3 asked you to look closely at how companies and researchers present information about the same technology. Did you notice any tension between how industry describes affect recognition and how your chosen sources describe it? If so, what do you think accounts for that tension?

Between Sessions

No new reading is required. Students should arrive to Session Two with Tasks 1, 2, 3, and the four reflection responses complete.

Session Two (80 minutes)

Opening: Jeopardy-Style Debrief of Session One — 20 minutes

Session Two opens with an instructor-led Jeopardy-style question board covering the work students completed in Session One — Tasks 2 and 3 and the four reflection questions. All questions are open-ended with no predetermined correct answers. The format invites students to surface what they actually found, rather than converging on a single right response. Categories map to the structure of the previous session's work: company research, source analysis, the AI document itself, Crawford, and a wildcard category for anything that genuinely caught students off guard.

This debrief serves as both a review and a bridge — consolidating Session One findings before students move into the more evaluative and comparative work of Tasks 4 and 5.

Task 4: Critical Engagement with the Document — 15 minutes

Drawing on everything read and investigated across both sessions, students identify the following three things from the Claude-authored document:

- **One claim they accept** — and explain why: What evidence, reasoning, or corroborating source supports their acceptance of this claim?
- **One claim they question or disagree with** — and explain why: What specifically gives them pause? Is it the evidence cited, the framing, the source, or the logic of the argument?
- **One claim they want to investigate further** — and explain why: What would they need to find out to be more confident in this claim? Where would they look?

Responses should be specific, naming the claim directly rather than describing it in general terms.

Task 5: Crawford and the Claude-Authored Document — 25 minutes

Students return to Chapter 5 of *Atlas of AI* and identify specific passages or arguments from Crawford that correspond to each of the three sections in the Claude-authored document. For each correspondence identified, students address the following:

- Does the Claude-authored document accurately represent Crawford's argument, or does it simplify, extend, or depart from it in some way?
- Did you discover any statements attributed to or drawn from Crawford that appear to be exaggerated, obscured, or paraphrased in a way that changes or softens their original meaning?
- Consider this carefully: How would Claude *know* what Crawford actually says? Crawford's book is a copyrighted text. Claude cannot quote it directly — it can only work from patterns in the text it was trained on, which may or may not include accurate representations of Crawford's precise arguments. What does this mean for how you should read any AI-generated document that claims to represent the ideas of a living author? What are the implications for research, attribution, and intellectual honesty?

This task is not designed to catch Claude in a mistake — it is designed to help students understand the structural limitations of AI-generated synthesis and to practice the habit of always returning to primary sources.

Closing — 10 minutes

Return to two items from the beginning of the lesson:

Task 0 responses. Ask students to re-read the two sentences they wrote before Session One about their expectations of an AI-written document. Ask: *"Would you write those same sentences now? What changed, and why?"* Take two or three responses aloud.

Question 4 from the pre-quiz. Return to the True/False item about emotional expression and genetics. Ask again for a show of hands. Note any shift from Session One. Ask: *"What specifically moved you, if you changed your answer?"*

This double return — to Task 0 and to the pre-quiz — gives students a concrete before-and-after marker of what shifted across the two sessions, and closes the metacognitive thread that has run through the entire lesson.

Submission

All written responses — Task 0 through Task 5, including the four Reflection prompts — are submitted via Brightspace at the end of Session Two or within 24 hours if additional time is needed.

Pacing Summary

Component	When	Time
Pre-class quiz	Before Session One	15–25 min
Task 0	Before Session One (out of class)	5 min
Opening discussion	Session One	5 min
Task 1: Read aloud (8 students × ~2 min) + brief debrief	Session One	20 min
Task 2: Company research	Session One	15 min
Task 3: Source analysis (2 sources)	Session One	25 min
Reflection R1–R4	Session One	15 min
Jeopardy-style debrief	Session Two	20 min
Task 4: Critical engagement	Session Two	15 min
Task 5: Crawford analysis	Session Two	25 min
Closing	Session Two	10 min
Total in-class time		**~155 min**

A Note on the AI-Authored Document

Students should be reminded — and the document itself states — that the primary reading was written by an AI. This is not incidental to the lesson; it is part of it. When students discover that a citation was corrected mid-process, that a paper was retracted, or that the AI's representation of Crawford differs from what Crawford actually wrote, they are not finding failures — they are practicing exactly the critical habits this lesson is designed to build. Instructors should feel free to name this explicitly during Session Two's debrief: *"The document is a primary source for today. Treat it like one."*

Lesson plan developed in consultation with Claude (claude.ai), March 2026.